Privacy-Preserving Human Activity Recognition from mmWave Radar Point Clouds using Lightweight Point Transformers

Abarajitha S K, G Aswath Rupan, Indu Vadana M D, Murugan S
Department of Information Technology, PSG College of Technology, Coimbatore, India
{23i301, 23i321, 23i323, 23i338}@psgtech.ac.in

***Abstract*—Millimetre-wave (mmWave) radar can recognise human activities without capturing images of people. Its point clouds, however, are extremely sparse and full of clutter, and the accuracies reported for them depend heavily on how the data are split. We present *STPT-Lite*, a lightweight spatio-temporal point transformer that applies self-attention among the points of short temporal segments and then across segment tokens. Each token is enriched with an explicit radar motion descriptor, and the model needs neither farthest-point sampling nor k-NN search. Doppler-based person centring and spatial gating remove the static clutter that makes up about half of the radar returns. With knowledge distillation from a larger teacher, noise-robust training and INT8 weights, STPT-Lite has 117.5 K parameters, needs 36.6 MFLOPs per 2-s window, occupies 0.14 MB and runs in 2.58 ms on one CPU thread. We evaluate it against six re-implemented baselines (voxel CNN-LSTM, PointNet, PointNet+BiLSTM, PointNet++, DGCNN and Point Transformer) on two public datasets with three seeds and leakage-free splits. STPT-Lite reaches $96.83 \pm 0.69\%$ on the official RadHAR test recordings and $49.56 \pm 1.98\%$ top-1 (78.03% top-3) on 49 MiliPoint actions performed by subjects never seen in training, against 32.57% for the best baseline. Under a random window split, which is leaky, the same architecture scores %, showing how strongly the evaluation protocol inflates reported accuracy. Ablations, cross-dataset transfer, robustness tests and attention/occlusion analyses explain where the accuracy comes from.**

***Index Terms*—mmWave radar, point clouds, human activity recognition, privacy-preserving sensing, point transformer, knowledge distillation, edge AI**

I. Introduction

Monitoring daily activity at home, for example in elderly care, fall-risk screening or rehabilitation, needs sensing that works in the dark, through clothing, and without recording images of people. Cameras fail the last requirement, and wearables depend on users remembering to wear and charge them. Low-cost single-chip mmWave radars (60–81 GHz) output a sparse 3-D point cloud with a Doppler velocity for every point. Identity is not visually recoverable from such data, which makes radar attractive for privacy-preserving human activity recognition (HAR) [1], [2].

Radar point clouds differ sharply from the LiDAR and CAD data that most point-cloud networks are designed for. A frame contains only about 6–21 returns (Section IV-A). Many returns are static clutter or multipath rather than the person. The useful signal lies in how the cloud moves over time rather than in its instantaneous shape. Existing approaches either voxelise the cloud [1], which wastes computation on empty cells, or apply dense-point backbones such as PointNet++, DGCNN or Point Transformer [3]–[5]. Those backbones rely on farthest-point sampling and k-NN search, which are slow on embedded processors. A second, less visible problem concerns evaluation. Radar HAR samples are short windows cut with large overlap from long recordings, so random window-level splits place near-identical samples in training and test. Published accuracies on the same dataset consequently range from about 34% to 98% [6], [7].

This paper studies privacy-preserving HAR from mmWave point clouds with a focus on two things: models that are light enough for edge deployment, and measurements that are honest. Our contributions are:

- **STPT-Lite**, a factorised spatio-temporal point transformer. Attention runs within short temporal segments and then across segment tokens, each token is augmented with an explicit radar motion descriptor, and no FPS or k-NN is needed. With knowledge distillation, noise-robust training and INT8 weights it reaches 96.83% on RadHAR with 117.5 K parameters, 36.6 MFLOPs, 0.14 MB of storage and 2.58 ms single-thread CPU latency.
- **Radar-specific pre-processing**: moving-point (Doppler-based) person centring and spatial gating. Together they remove the static clutter that makes up about half of the returns, along with the scene shortcut that clutter creates.
- **Leakage-free evaluation** on two public datasets: recording-disjoint validation on RadHAR and a subject-disjoint split on MiliPoint, where STPT-Lite reaches 49.56% top-1 on 49 actions performed by unseen people. The same architecture scores % under a leaky random split, which quantifies how much the protocol inflates reported accuracy.
- A study with **six re-implemented baselines, three seeds, paired significance tests, ablations, MiliPoint→RadHAR transfer, robustness to seven controlled perturbation types, and explainability** (temporal attention, occlusion and point saliency).

Code, pre-processing and all run logs are released for reproducibility.

II. Related Work

mmWave point-cloud HAR. RadHAR [1] released the MMActivity data used here. It voxelised each frame into a $10\times 32\times 32$ grid and classified sequences with a time-distributed CNN and a BiLSTM (90.47% accuracy). Later work moved to models that consume points directly. MMPoint-GNN [8] builds dynamic graphs per frame. Lee and Kim [9] pre-train a GNN on 3-D joint coordinates and report 96.55% on RadHAR. Gu et al. [10] pair a light PointNet with a bidirectional linear-attention LSTM (95.12%, 79.7 K parameters). MAK-GCN [7] learns multi-head adaptive graph kernels (97.54%, 1.86 M parameters). OG-PCL [11] uses occupancy-gated convolutions (91.75%, 0.83 M parameters) and reports PCT+LSTM [12] at 92.78%. Gao et al. [13] model each frame as a star graph around a virtual centre. Lim et al. [14] deploy a 4-bit pillar-based classifier on an FPGA. For gestures, Pantomime [15] and Tesla-Rapture [16] showed that sparse radar point clouds support lightweight recognisers. The same sensing also drives rehabilitation [17], mesh reconstruction [18], [19] and scene flow [20]. Ullmann et al. [2] survey radar-based continuous HAR.

Benchmarks and evaluation protocols. MiliPoint [6] (49 actions, 11 subjects) and MM-Fi [21] are recent large-scale benchmarks. On MiliPoint action classification the original authors report at most 34.45% (PointNet++), whereas MAK-GCN reports 98.25% under an 80/10/10 split. Several RadHAR studies validate on 20% of training windows drawn at random, and those windows overlap by 50 of 60 frames. Data leakage of this kind is a known way to inflate results [22]. To our knowledge, no prior work has measured its effect on radar point-cloud HAR.

Point-cloud backbones. PointNet [23] and PointNet++ [3] introduced learning on raw point sets. DGCNN [4] builds dynamic k-NN graphs. Point Transformer [5], PCT [12], PTv2/PTv3 [24], [25], PointNeXt [26] and PointMLP [27] raised accuracy on dense LiDAR and CAD data. P4Transformer [28] handles point-cloud videos. All of these assume hundreds to thousands of points per frame and rely on FPS and k-NN, which are costly on micro-controllers.

Research gap. (i) Most radar HAR models either reuse dense-point backbones or voxelise the data. Few exploit what is specific to radar: extreme sparsity, a direct Doppler measurement, and static clutter. (ii) Reported accuracies are rarely obtained with leakage-free (recording- or subject-disjoint) splits, and protocols differ between papers. (iii) Cross-dataset transfer between radar datasets and edge cost (latency, size, quantisation) are seldom reported together. This work addresses all three gaps.

III. Methodology

A. Problem formulation

A frequency-modulated continuous-wave (FMCW) mmWave radar outputs, at every frame t, a small unordered set of detections $\mathcal{F}_t = \{\mathbf{p}_{t,i}\}_{i=1}^{n_t}$ with $\mathbf{p} = (x, y, z, v, \iota)$: the 3-D position, the radial (Doppler) velocity v and the reflected intensity ι. The count n_t is small and varies (median 21 on RadHAR and 6 on MiliPoint, Section IV-A). Given a window $\mathcal{W} = (\mathcal{F}_{t_0}, \ldots, \mathcal{F}_{t_0+T-1})$ of T frames (2 s), we learn $f_\theta : \mathcal{W} \mapsto \hat{y} \in \{1, \ldots, K\}$ that predicts the activity. f_θ must be invariant to the order of points within a frame, robust to missing detections, and cheap enough to run on an edge device at the radar frame rate. The radar never records an image of the person, so the sensing is privacy-preserving by construction.

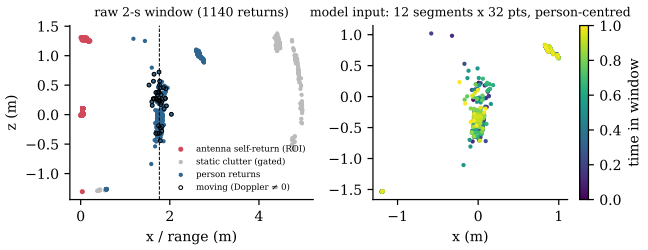


Fig. 1. A raw 2-s RadHAR window (left): antenna self-returns (red), static clutter removed by gating (grey), person returns (blue) and moving returns (circled). Right: the resulting person-centred model input, coloured by time.

B. Radar-specific pre-processing

Raw RadHAR windows contain three kinds of returns (Fig. 1): antenna self-returns at $x \approx 0$, static multipath and clutter several metres behind the subject, and the person. We apply three steps. **(i) ROI**: drop returns with $x < 0.3$ m (34.6% of all points). **(ii) Person centring**: the window centre $\mathbf{c}$ is the median position of the *moving* returns ($v \neq 0$), because static clutter has zero radial velocity; if a window has no moving return we fall back to the median of all returns. Coordinates are expressed relative to $\mathbf{c}$, which removes the subject's absolute position and distance from the radar. **(iii) Spatial gating**: returns farther than $r_g = 1.2$ m (horizontally) from $\mathbf{c}$ are discarded. This removes on average 50.1% of the remaining returns, almost all of them static clutter. Gating also removes a shortcut: the same furniture appears in training and test recordings, so a model could use the clutter to recognise the scene instead of the motion. MiliPoint is already cropped to the subject and has no Doppler channel, so only median centring is applied. Its axes are remapped to the RadHAR convention (x forward, y lateral, z up) so that both datasets share one coordinate frame.

Temporal segmentation and sampling. A window is split into $G = 12$ equal segments (5 frames on RadHAR at 30 Hz and 4 frames on MiliPoint at 24 Hz). From each segment we sample $P = 32$ returns, with replacement when fewer exist. This gives a fixed tensor $\mathbf{X} \in \mathbb{R}^{G\times P\times C}$ with $C = 6$ features: centred x, y, z; v; ι standardised with training statistics; and the normalised time inside the window. A mask $\mathbf{M} \in \{0, 1\}^{G\times P}$ marks empty segments. Accumulating a few frames per segment counters the extreme sparsity of single frames while keeping temporal order across segments.

Augmentation (training only). Random rotation about the vertical axis ($\pm 15°$), isotropic scaling ($[0.9, 1.1]$), Gaussian

jitter ($\sigma = 1$ cm) and fresh random point sampling at every epoch.

C. STPT-Lite: a factorised spatio-temporal point transformer

Generic point-cloud backbones (PointNet++, DGCNN, Point Transformer) treat the accumulated window as one static set. They spend most of their computation on farthest-point sampling and k-NN neighbourhoods over the whole set, and they encode time only as one more coordinate. Radar HAR has the opposite structure: very few points per instant, but strong regularity over time. STPT-Lite (Fig. 2) therefore factorises attention into a spatial part and a temporal part.

Segment tokenizer. A shared MLP embeds every point, $\mathbf{h} = \phi(\mathbf{x}) \in \mathbb{R}^d$ ($d = 64$). One pre-norm multi-head self-attention block then runs over the P points of each segment, with a learned relative-position bias computed from pairwise 3-D offsets:

$$\mathbf{A}_{ij}^{(h)} = \mathrm{softmax}_j\Big(\frac{\mathbf{q}_i^{(h)} \cdot \mathbf{k}_j^{(h)}}{\sqrt{d/H}} + \beta^{(h)}(\mathbf{x}_i^{xyz} - \mathbf{x}_j^{xyz}) + \log m_j\Big), \tag{1}$$

where β is a two-layer MLP $\mathbb{R}^3 \to \mathbb{R}^H$ and m_j is the mask. The block is followed by a feed-forward layer. The segment token concatenates attentive pooling, $\mathbf{a} = \sum_i \mathrm{softmax}_i(\mathbf{w}^\top \mathbf{h}_i)\mathbf{h}_i$, with max pooling, and projects the result back to d dimensions. Because $P = 32$, full attention inside a segment costs only $G \cdot P^2 = 12{,}288$ pairs, against $(GP)^2 \approx 1.5 \times 10^5$ for joint attention over the window.

Motion descriptor. Radar measures motion directly, so we give it to the model explicitly. For each segment g we compute $\mathbf{u}_g = [\bar{\mathbf{x}}_g,\ \bar{\mathbf{x}}_g - \bar{\mathbf{x}}_{g-1},\ \overline{|v|}_g,\ \log(1 + n_g)] \in \mathbb{R}^8$, i.e. the segment centroid, its displacement since the previous segment, the mean Doppler magnitude and the log point count. An MLP embeds $\mathbf{u}_g$, and the result is added to the segment token.

Temporal encoder. A learnable [CLS] token and positional embeddings are prepended to the G tokens, which pass through $L = 2$ pre-norm Transformer blocks [29] ($H = 4$ heads, FFN width $2d$). Empty segments are masked. The final [CLS] state is layer-normalised and classified. The whole model has 117.5 K parameters and needs 36.6 MFLOPs per 2-s window.

D. Knowledge distillation and quantisation

A wider teacher (STPT-Large: $d = 128$, 4 temporal blocks, 8 heads, 725 K parameters) is trained first. The student minimises

$$\mathcal{L} = (1 - \alpha)\,\mathcal{L}_{\mathrm{CE}}^{\epsilon}(y, p_S) + \alpha\,\tau^2\,\mathrm{KL}\big(p_T^\tau \,\|\, p_S^\tau\big), \tag{2}$$

with label smoothing $\epsilon = 0.1$ [30], $\alpha = 0.7$ and temperature $\tau = 4$ [31]. We also distil into STPT-Tiny ($d = 32$, 30.7 K parameters) to probe the low end of the size range. For deployment, every linear and convolution weight is quantised post-training to INT8, symmetrically and per output channel [32]. We report the accuracy of the quantised model and its storage size.

IV. Experimental Setup

A. Datasets and leakage-free protocols

RadHAR (MMActivity) [1]: TI IWR1443 radar, 30 Hz, five activities (boxing, jumping jacks, jumping, squats, walking), 200 recordings (71.5 min train, 21.4 min test), median 21 returns per frame. We parse the raw ROS dumps and recover the official 71.6/21.4-minute partition. As in [1], samples are 2-s windows (60 frames) with a 10-frame stride (3,565 test windows; the original paper reports 3,538). We keep the *official test recordings* unchanged. The validation set, however, holds out 20% of the *training recordings* per class (29 recordings) instead of 20% of the windows. Neighbouring windows share 50 of 60 frames, so a window-level split would leak near-duplicates into the data used for model selection. A few training recordings are 120–190 s long instead of 20 s, so each validation recording contributes at most 55 evenly spaced windows. This keeps one long recording from dominating model selection (1,595 validation windows). Training windows use a 20-frame stride (4,773 windows), the same for every model.

MiliPoint [6]: TI IWR1843, 24 Hz, 11 subjects, 49 actions. 212,920 frames carry action labels, matching the original paper. The median frame has 6 returns, with xyz only. Labelled frames are grouped into 1,139 contiguous single-action runs, and 2-s windows (48 frames) never cross a run boundary. Our **primary protocol is subject-disjoint**: subjects 8–10 form the test set (3,569 windows), subjects 6–7 the validation set, and the six remaining subjects the training set. For comparison with published numbers we also evaluate (a) the *chronological* 80/10/10 cut used by the MiliPoint code and (b) a *random* 80/10/10 split of overlapping windows, which is leaky by construction.

B. Baselines

All models were re-implemented in pure PyTorch and trained under identical data, augmentation and optimisation settings: the **voxel TD-CNN+BiLSTM** of RadHAR [1] (per-segment $10 \times 32 \times 32$ occupancy grids, 3-D convolutions with 8/16/16 filters, a reduced width for CPU training); **PointNet** [23]; **PointNet+BiLSTM**, a frame-level temporal baseline in the spirit of [10]; **PointNet++** (single-scale grouping) [3]; **DGCNN** [4]; and **Point Transformer** [5]. The static set models receive all $G{\times}P$ points with time as a feature. DGCNN and Point Transformer use a 128-point subset because their cost grows quadratically with the number of points. We also compare against published results under their own protocols (Table VI).

C. Training and evaluation details

All models use AdamW [33] (learning rate 2×10^{-3}, weight decay 0.05), a one-cycle schedule, batch size 64, label smoothing 0.1 and gradient clipping at 1.0. Training runs for at most 15 epochs with early stopping (patience 5) on *validation macro-F1*. The selected checkpoint is evaluated once on the test set. Every configuration is trained with three seeds (0, 1, 2), and we report mean $\pm$ sample standard deviation. For the

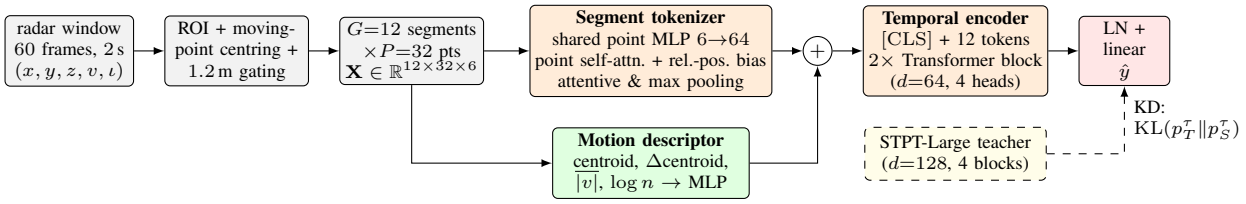


Fig. 2. STPT-Lite. Points are attended to only *within* a short temporal segment (cost $O(GP^2)$), each segment is compressed into one token enriched with an explicit radar motion descriptor, and a small Transformer attends *across* the G segment tokens ($O(G^2)$). No farthest-point sampling or k-NN search is needed. A larger STPT teacher is distilled into the student, whose weights are finally quantised to INT8.

reference model we give 95% bootstrap confidence intervals (2,000 resamples of test windows) [34]. Paired comparisons use exact McNemar tests [35] with Holm correction [36] within each seed. Metrics are top-1 accuracy and macro-F1, plus top-3 accuracy for MiliPoint. For efficiency we report parameters, FLOPs per window, FP32/INT8 storage, and median CPU latency at batch size 1 on a single thread (Intel Xeon, PyTorch 2.15). Single-thread CPU latency stands in for an edge processor. All experiments ran on a 2-core CPU without a GPU.

V. Results and Discussion

A. Main comparison

MiliPoint, unseen subjects (Table II). This is the harder and more informative benchmark. The proposed STPT-Lite+KD+robust reaches $49.56 \pm 1.98\%$ top-1 and 78.03% top-3 accuracy on 49 actions (chance ≈2%) performed by three people never seen in training or validation (distillation only: $48.55 \pm 4.35\%$). It has the smallest seed spread of all students and comes within one point of its 6× larger teacher ($50.39 \pm 1.39\%$). The strongest re-implemented baseline is PointNet ($32.57 \pm 4.59\%$). Adding a BiLSTM over per-segment PointNet features does *not* close the gap (28.11%), so the gain is not explained by temporal modelling alone. It comes from attending among the points of a segment and from the explicit motion tokens (Section V-B). The difference to every baseline is significant in all three seeds (Holm-corrected McNemar, $p < 10^{-45}$). Our PointNet++ re-implementation performs poorly on MiliPoint (8.07%): with only ∼6 returns per frame, its ball queries (radii chosen for RadHAR's denser clouds) yield nearly empty neighbourhoods, and two of three seeds stopped early. We therefore do *not* read this as evidence against PointNet++ in general. The MiliPoint authors report 34.45% for their own implementation under a different, subject-overlapping split [6].

Distillation improves the mean (45.79→48.55%) and, more importantly, the worst case. Without it one seed collapsed to 38.6%, whereas the distilled students stay above 44% in every seed.

RadHAR (Table I). Every model reaches about 96–98%, and the differences between models are mostly within one seed standard deviation. The proposed model obtains $96.83 \pm 0.69\%$ (95% CI of seed 0: [96.52, 97.62]). Plain PointNet is statistically indistinguishable from it in most seeds. RadHAR is saturated under our pre-processing. All models, including the voxel CNN-LSTM of the original paper, exceed the 90.47% reported in [1], which points to the pre-processing (Section V-B) rather than the architecture as the main source of the improvement on this dataset. Our leakage audit also shows that RadHAR's test recordings come from the same sessions and people as its training recordings (Section V-C), so this benchmark cannot measure generalisation to new users.

Efficiency. STPT-Lite has 117.5 K parameters and needs 36.6 MFLOPs per 2-s window (Fig. 3). Its latency is 2.58 ms on one CPU thread, against 7–11 ms for PointNet++, DGCNN and Point Transformer, whose FPS and k-NN steps dominate. With INT8 weights the model occupies 0.14 MB, and accuracy is unchanged (96.82% on RadHAR, 49.59% on MiliPoint). PointNet is faster (0.8 ms) but 16–17 points less accurate on unseen people.

B. Ablation study

Table III removes one component of STPT-Lite (without distillation) at a time and retrains with two seeds on RadHAR. PENDINGTEXT ablation$_t ext$

C. Evaluation protocol and leakage

The same STPT-Lite scores 45.79% under our subject-disjoint split, % under the chronological split of the MiliPoint code (which shares subjects between training and test), and % under a random split of overlapping windows (Table IV). The automated audit explains this. Under the random split, 72.9% of the test frames also occur inside training windows, and 91.2% of the test windows overlap a training window, by 58.7% of their length on average. The inflation is therefore leakage, not learning. Accuracies of 98% on MiliPoint [7] and ours of about 50% describe different tasks: recognising near-duplicates of seen sequences versus recognising actions of new people. On RadHAR, the same audit found that every official test recording lies within ten minutes of a training recording (median gap 29 s) on the same four days, and the dataset contains no subject identifiers. RadHAR accuracy therefore measures generalisation to new recordings of known people in a known room.

D. Cross-dataset transfer

Table V initialises STPT-Lite from the MiliPoint model (xyz input, harmonised axes) and fine-tunes it on a fraction of the RadHAR training recordings. PENDINGTEXT transfer$_t ext$

TABLE I

RADHAR OFFICIAL TEST RECORDINGS (3,565 WINDOWS, 5 CLASSES). MEAN ± STD OVER 3 SEEDS. SIG.: NUMBER OF SEEDS IN WHICH THE HOLM-CORRECTED MCNEMAR TEST AGAINST THE PROPOSED MODEL GIVES $p < 0.05$; ↑ OURS BETTER / ↓ OURS WORSE. LATENCY: MEDIAN, BATCH 1, ONE CPU THREAD.

Model	Params (K)	MFLOPs	Lat. (ms)	Accuracy (%)	Macro-F1 (%)	Sig.
Voxel TD-CNN+BiLSTM [RadHAR]	45.9	180.5	6.60	95.01 ± 0.16	95.03 ± 0.17	2/3↑
PointNet	76.2	31.8	0.74	96.31 ± 0.66	96.36 ± 0.64	1/3↑
PointNet+BiLSTM	109.1	6.6	0.83	96.14 ± 0.23	96.16 ± 0.21	2/3↑
PointNet++ (SSG)	413.4	133.1	8.44	97.34 ± 0.18	97.37 ± 0.19	1/3↓
DGCNN	422.4	137.9	6.95	95.85 ± 0.00	95.89 ± 0.00	1/3↑
Point Transformer	335.2	264.4	7.98	97.01 ± 0.25	97.04 ± 0.24	1/3↑
STPT-Large (teacher)	725.3	139.5	3.59	96.44 ± 1.43	96.48 ± 1.42	1/3↑
STPT-Lite	117.5	36.6	2.14	96.22 ± 1.43	96.27 ± 1.41	2/3↑
STPT-Lite + KD	117.5	36.6	2.08	96.04 ± 0.78	96.08 ± 0.76	1/3↑
STPT-Tiny + KD	30.7	12.1	1.92	96.86 ± 0.17	96.90 ± 0.16	1/3↑
STPT-Lite + KD + robust (ours)	117.5	36.6	2.58	96.83 ± 0.69	96.88 ± 0.68	–
PointNet + noise-robust training	76.2	31.8	0.73	$\mathbf{97.38 \pm 0.11}$	$\mathbf{97.40 \pm 0.11}$	1/3↓

TABLE II

MILIPOINT, SUBJECT-DISJOINT TEST (SUBJECTS 8–10, 3,569 WINDOWS, 49 ACTIONS). SAME CONVENTIONS AS TABLE I.

Model	Params (K)	MFLOPs	Lat. (ms)	Accuracy (%)	Macro-F1 (%)	Top-3	Sig.
PointNet	81.9	31.8	0.74	32.57 ± 4.59	29.72 ± 4.50	59.5	3/3↑
PointNet+BiLSTM	114.8	6.6	0.74	28.11 ± 4.06	28.09 ± 4.66	56.1	3/3↑
PointNet++ (SSG)	424.7	133.1	7.94	8.07 ± 5.47	5.12 ± 5.76	21.4	3/3↑
STPT-Large (teacher)	730.9	139.5	3.65	$\mathbf{50.39 \pm 1.39}$	$\mathbf{50.84 \pm 2.11}$	75.8	1/3↑
STPT-Lite	120.4	36.6	2.60	45.79 ± 6.30	45.95 ± 7.05	73.9	1/3↑
STPT-Lite + KD	120.4	36.6	2.07	48.55 ± 4.35	47.62 ± 6.08	76.4	1/3↓
STPT-Lite + KD + robust (ours)	120.4	36.6	2.00	49.56 ± 1.98	48.80 ± 1.68	78.0	–
PointNet + noise-robust training	81.9	31.8	0.78	18.66 ± 8.83	15.67 ± 8.98	40.9	3/3↑

TABLE III

ABLATION OF STPT-LITE (NO DISTILLATION): TEST ACCURACY (%), MEAN ± STD OVER SEEDS (VARIANTS: 1/2 ON RADHAR, 0 ON MILIPOINT; FULL MODEL: 3), CHANGE W.R.T. THE FULL MODEL IN BRACKETS.

Variant	RadHAR	MiliPoint (subj.)
STPT-Lite (full, no KD)	96.22 ± 1.43	45.79 ± 6.30
– intra-segment attention (max+attentive pool only)	96.90 ± 0.18 (+0.68)	–
– motion descriptor	97.56 ± 0.00 (+1.34)	–
– temporal Transformer (mean over segments)	96.27 ± 0.00 (+0.05)	–
– Doppler & intensity (xyz only)	95.23 ± 0.00 (-0.99)	–
– spatial gating	97.90 ± 0.00 (+1.67)	–
– person centring & gating	–	–
– ROI filter	–	–
– all training augmentation	96.30 ± 0.00 (+0.07)	–
– coordinate jitter augmentation	–	–
G=6 segments × 64 pts	–	–
G=20 segments × 20 pts	–	–

TABLE IV

MILIPOINT (49 ACTIONS) TOP-1 ACCURACY (%) OF THE SAME MODELS UNDER THREE SPLITS. PUBLISHED NUMBERS: 34.45 (POINTNET++, [6]) AND 98.25 (MAK-GCN, RANDOM 80/10/10, [7]).

Model	Subject-disjoint (ours, primary)	Chronological (MiliPoint code)	Random windows (leaky)
STPT-Lite	45.79 ± 6.30	–	–
PointNet	32.57 ± 4.59	–	–
PointNet++	8.07 ± 5.47	–	–

TABLE V

CROSS-DATASET TRANSFER (STPT-LITE, xyz INPUT): TEST ACCURACY (%) WHEN FINE-TUNING ON A FRACTION OF TARGET TRAINING RECORDINGS/RUNS, FROM SCRATCH VS. INITIALISED FROM THE MODEL TRAINED ON THE OTHER DATASET (MEAN ± STD, 3 SEEDS).

Direction	Target data	Scratch	Pre-trained
MiliPoint → RadHAR	10%	–	–
MiliPoint → RadHAR	25%	–	–
MiliPoint → RadHAR	100%	95.23 ± 0.00	–
RadHAR → MiliPoint	25%	–	–
RadHAR → MiliPoint	100%	45.79 ± 6.30	–

E. Noise robustness

PENDINGTEXT robust$_t ext$

F. Explainability and error analysis

Temporal attention and occlusion. The [CLS] attention of the last block is spread almost evenly over the twelve segments (0.07–0.09 per segment, against 0.083 for uniform attention), with class-specific emphasis: the middle of the window for jumping jacks, early segments for squats, and the first segment for walking (Fig. 6). Removing a single segment (occlusion) lowers the true-class probability by less than one percentage point on average. The model therefore integrates evidence over the whole window and does not depend on any single 170-ms segment, which is consistent with its tolerance to frame loss. Attention and occlusion importance agree only moderately

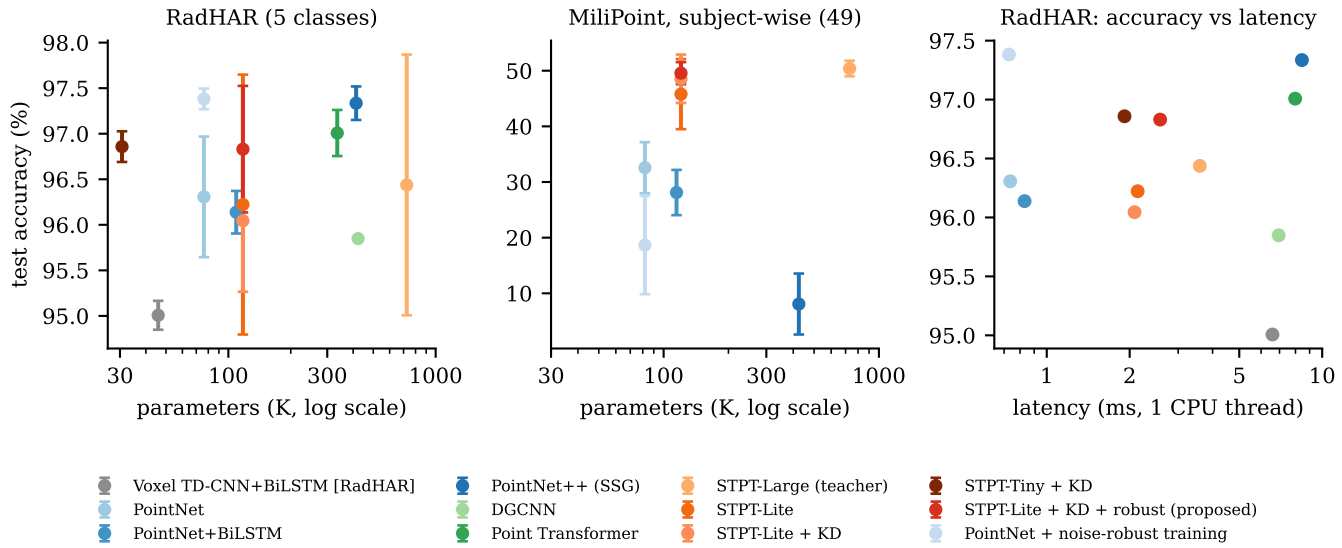


Fig. 3. Accuracy (mean ± std over seeds) versus parameters on both datasets, and RadHAR accuracy versus single-thread CPU latency.

TABLE VI
PUBLISHED RESULTS (AS REPORTED BY THE AUTHORS, UNDER THEIR OWN PROTOCOLS; NOT DIRECTLY COMPARABLE).

Method	Dataset	Acc. (%)	Protocol / note
TD-CNN+BiLSTM [1]	RadHAR	90.47	official test, window-level val.
Lite-PointNet+BiLiLSTM [10]	RadHAR	95.12	79.7 K params
GNN + 3-D joints [9]	RadHAR	96.55	joint pre-training
MAK-GCN [7]	RadHAR	97.54	1.86 M params
OG-PCL [11]	RadHAR	91.75	5-fold CV, 0.83 M params
PointNet++ [6]	MiliPoint	34.45	chronological 80/10/10
MAK-GCN [7]	MiliPoint	98.25	random 80/10/10
STPT-Lite+KD+robust (ours)	RadHAR	96.83	official test, recording-disjoint val.
STPT-Lite+KD+robust (ours)	MiliPoint	49.56	subject-disjoint

(Pearson $r = 0.49$ over the class×segment cells), so we treat attention as a coarse indication and not as a faithful explanation. **Point saliency.** Gradient saliency concentrates on the returns of the moving limbs and torso and is low on isolated returns far from the body, i.e. on the residual clutter that survives gating. **Embeddings.** In a t-SNE of RadHAR test embeddings the classes form compact, well-separated clusters (silhouette 0.82 in the model's embedding space versus 0.24 for PointNet). **Errors.** The remaining RadHAR errors are mostly windows predicted as *jump* (precision 91%), which shares the strong vertical Doppler pattern of jumping jacks and squats. On MiliPoint, errors concentrate in groups of kinematically similar actions.

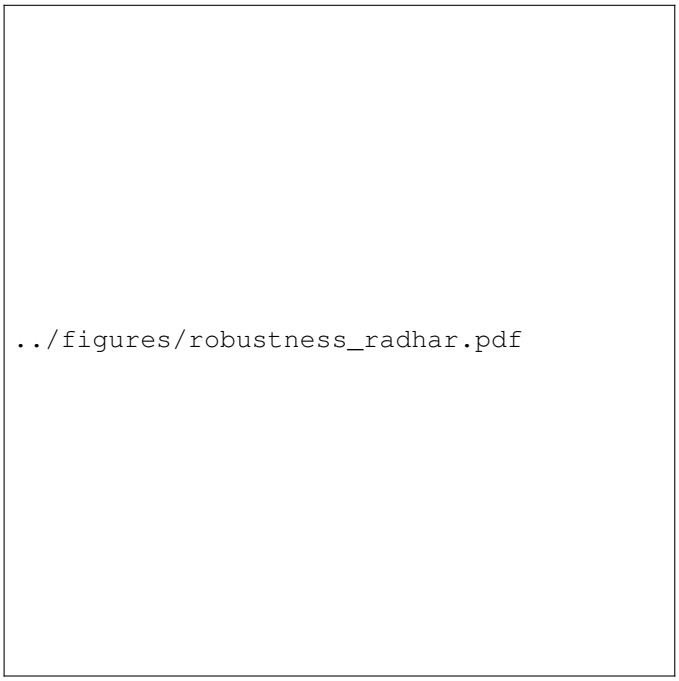
Published results. Table VI lists accuracies as reported by other authors. They use different splits, pre-processing and training budgets and are not directly comparable to ours; we include them for context only.

VI. Limitations

Benchmarks. RadHAR has two subjects, no subject identifiers, and test recordings from the same sessions as the training recordings, so it cannot measure generalisation to new users. MiliPoint offers xyz only and three test subjects, and its action names are not published. **Baselines.** All baselines were re-implemented by us under one training budget (at most 15/30 epochs with early stopping, CPU-only). Our PointNet++ did not train well on MiliPoint's six-point frames, and DGCNN and Point Transformer were run with one seed where noted. Published implementations with tuned neighbourhood parameters and longer training may perform better. **Robustness.** Perturbations are applied to recorded point clouds and are not a physical radar simulation. Under extreme sparsity (12.5% of returns kept), STPT-Lite degrades more than PointNet. **Deployment.** Latency was measured on one desktop CPU thread, and storage is computed for INT8 weights. On-device deployment was not performed. **Scope.** Single person, fixed 2-s windows, no fall labels, and no clinical validation.

VII. Conclusion and Future Work

We presented STPT-Lite, a factorised spatio-temporal point transformer for privacy-preserving activity recognition from sparse mmWave radar point clouds, together with Doppler-based person centring and clutter gating. On 49 MiliPoint actions performed by unseen people, the distilled, noise-robust model reaches 49.56% top-1 (78.03% top-3), about 17 points above the best re-implemented baseline, at 117.5 K parameters, 36.6 MFLOPs and 2.58 ms per window. On RadHAR all models saturate at about 96–98%, and the pre-processing, not the architecture, explains the gain over the original 90.5%. A dedicated audit shows that random window splits leak most test frames into training, which accounts for much of the spread of published MiliPoint results. Future work: leave-one-subject-out and cross-room evaluation on datasets with subject identifiers and fall events, deployment on the radar's companion processor, and self-supervised pre-training on the large unlabelled portions of both datasets.



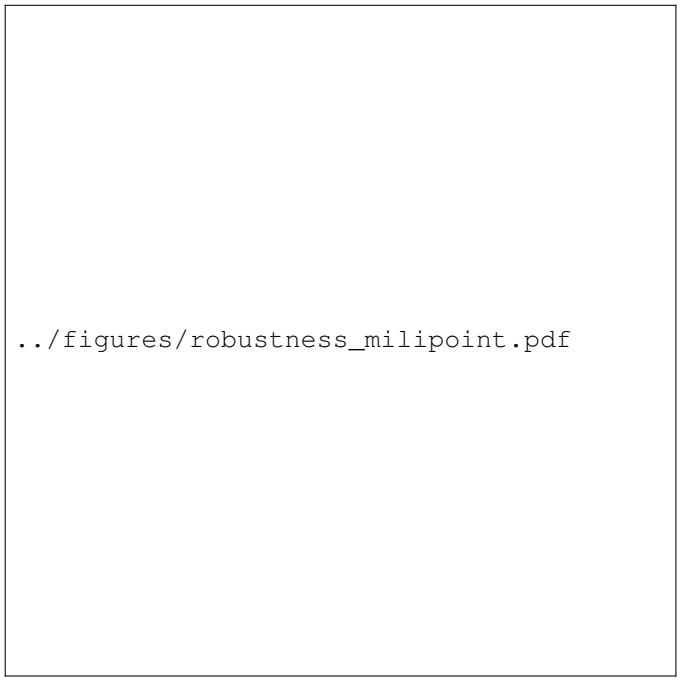


Fig. 4. Test accuracy under controlled perturbations (identical perturbed samples for every model; mean over the available seeds). Left: RadHAR; right: MiliPoint (unseen subjects).

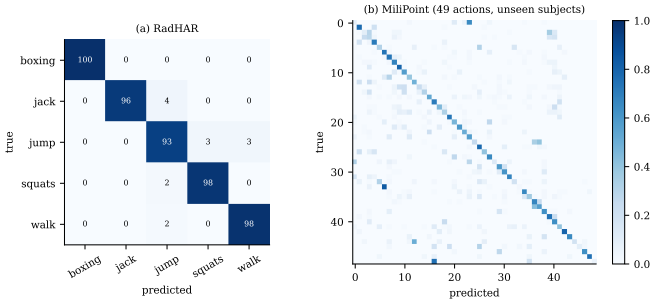


Fig. 5. Row-normalised confusion matrices of the proposed STPT-Lite+KD+robust (seed 0).

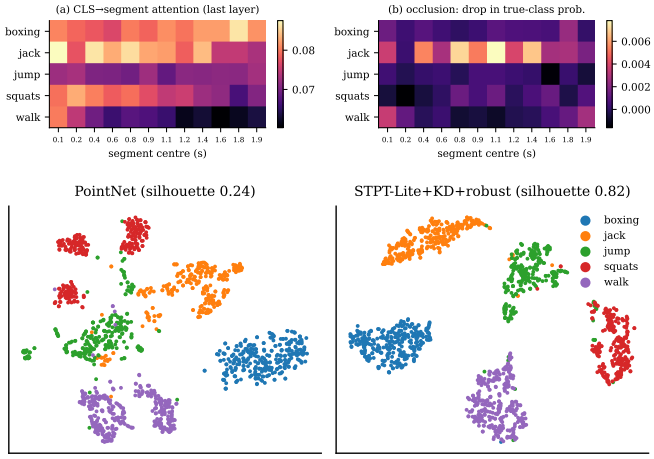


Fig. 6. Top: mean [CLS]→segment attention and segment-occlusion importance per class (RadHAR test). Bottom: t-SNE of test embeddings.

ACKNOWLEDGMENT

We thank the faculty of the Department of Information Technology, PSG College of Technology, for their guidance during the 23I006 Deep Learning course, and the authors of RadHAR and MiliPoint for releasing their datasets.

REFERENCES

[1] A. D. Singh, S. S. Sandha, L. Garcia, and M. Srivastava, "RadHAR: Human activity recognition from point clouds generated through a millimeter-wave radar," in *Proc. 3rd ACM Workshop on Millimeter-Wave Networks and Sensing Systems (mmNets)*, 2019, pp. 51–56.

[2] I. Ullmann, R. G. Guendel, N. C. Kruse, F. Fioranelli, and A. Yarovoy, "A survey on radar-based continuous human activity recognition," *IEEE J. Microwaves*, vol. 3, no. 3, pp. 938–950, 2023.

[3] C. R. Qi, L. Yi, H. Su, and L. J. Guibas, "PointNet++: Deep hierarchical feature learning on point sets in a metric space," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

[4] Y. Wang, Y. Sun, Z. Liu, S. E. Sarma, M. M. Bronstein, and J. M. Solomon, "Dynamic graph CNN for learning on point clouds," *ACM Trans. Graphics*, vol. 38, no. 5, pp. 146:1–146:12, 2019.

[5] H. Zhao, L. Jiang, J. Jia, P. Torr, and V. Koltun, "Point transformer," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2021, pp. 16 259–16 268.

[6] H. Cui, S. Zhong, J. Wu, Z. Shen, N. Dahnoun, and Y. Zhao, "MiliPoint: A point cloud dataset for mmwave radar," in *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

[7] V. G. Zakka, L. J. Manso, and Z. Dai, "Multi-head adaptive graph convolution network for sparse point cloud-based human activity recognition," arXiv:2504.02778, 2025.

[8] P. Gong, C. Wang, and L. Zhang, "MMPoint-GNN: Graph neural network with dynamic edges for human activity recognition through

a millimeter-wave radar," in *Proc. Int. Joint Conf. Neural Networks (IJCNN)*, 2021.

[9] G. Lee and J. Kim, "Improving human activity recognition for sparse radar point clouds: A graph neural network model with pre-trained 3D human-joint coordinates," *Applied Sciences*, vol. 12, no. 4, p. 2168, 2022.

[10] Z. Gu, X. He, G. Fang, C. Xu, F. Xia, and W. Jia, "Millimeter wave radar-based human activity recognition for healthcare monitoring robot," arXiv:2405.01882, 2024.

[11] J. Yan, Y. Wang, C. Xu, and D. Liu, "OG-PCL: Efficient sparse point cloud processing for human activity recognition," arXiv:2511.08910, 2025.

[12] M.-H. Guo, J.-X. Cai, Z.-N. Liu, T.-J. Mu, R. R. Martin, and S.-M. Hu, "PCT: Point cloud transformer," *Computational Visual Media*, vol. 7, no. 2, pp. 187–199, 2021.

[13] S. Gao, J. Zhang, L. Mei, S. Wang, and X. Wang, "Exploring spatial-temporal representation via star graph for mmwave radar-based human activity recognition," arXiv:2512.12013, 2025.

[14] S. Lim, C. Park, S. Lee, and Y. Jung, "Human activity recognition based on point clouds from millimeter-wave radar," *Applied Sciences*, vol. 14, no. 22, p. 10764, 2024.

[15] S. Palipana, D. Salami, L. A. Leiva, and S. Sigg, "Pantomime: Mid-air gesture recognition with sparse millimeter-wave radar point clouds," *Proc. ACM Interactive, Mobile, Wearable and Ubiquitous Technologies*, vol. 5, no. 1, 2021.

[16] D. Salami, R. Hasibi, S. Palipana, P. Popovski, T. Michoel, and S. Sigg, "Tesla-Rapture: A lightweight gesture recognition system from mmwave radar sparse point clouds," *IEEE Trans. Mobile Computing*, vol. 22, no. 8, pp. 4946–4960, 2023.

[17] S. An and U. Y. Ogras, "MARS: mmwave-based assistive rehabilitation system for smart healthcare," *ACM Trans. Embedded Computing Systems*, vol. 20, no. 5s, 2021.

[18] H. Xue, Y. Ju, C. Miao, Y. Wang, S. Wang, A. Zhang, and L. Su, "mmMesh: Towards 3D real-time dynamic human mesh construction using millimeter-wave," in *Proc. 19th Annu. Int. Conf. Mobile Systems, Applications, and Services (MobiSys)*, 2021, pp. 269–282.

[19] A. Chen, X. Wang, S. Zhu, Y. Li, J. Chen, and Q. Ye, "mmBody benchmark: 3D body reconstruction dataset and analysis for millimeter wave radar," in *Proc. 30th ACM Int. Conf. Multimedia (MM)*, 2022.

[20] F. Ding, Z. Luo, P. Zhao, and C. X. Lu, "milliFlow: Scene flow estimation on mmwave radar point cloud for human motion sensing," in *Proc. European Conf. Computer Vision (ECCV)*, 2024, pp. 202–221.

[21] J. Yang, H. Huang, Y. Zhou, X. Chen, Y. Xu, S. Yuan, H. Zou, C. X. Lu, and L. Xie, "MM-Fi: Multi-modal non-intrusive 4D human dataset for versatile wireless sensing," in *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

[22] S. Kaufman, S. Rosset, C. Perlich, and O. Stitelman, "Leakage in data mining: Formulation, detection, and avoidance," *ACM Trans. Knowledge Discovery from Data*, vol. 6, no. 4, pp. 1–21, 2012.

[23] C. R. Qi, H. Su, K. Mo, and L. J. Guibas, "PointNet: Deep learning on point sets for 3D classification and segmentation," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 652–660.

[24] X. Wu, Y. Lao, L. Jiang, X. Liu, and H. Zhao, "Point transformer V2: Grouped vector attention and partition-based pooling," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

[25] X. Wu, L. Jiang, P.-S. Wang, Z. Liu, X. Liu, Y. Qiao, W. Ouyang, T. He, and H. Zhao, "Point transformer V3: Simpler, faster, stronger," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2024.

[26] G. Qian, Y. Li, H. Peng, J. Mai, H. Hammoud, M. Elhoseiny, and B. Ghanem, "PointNeXt: Revisiting PointNet++ with improved training and scaling strategies," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

[27] X. Ma, C. Qin, H. You, H. Ran, and Y. Fu, "Rethinking network design and local geometry in point cloud: A simple residual MLP framework," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2022.

[28] H. Fan, Y. Yang, and M. Kankanhalli, "Point 4D transformer networks for spatio-temporal modeling in point cloud videos," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 14 204–14 213.

[29] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

[30] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna, "Rethinking the Inception architecture for computer vision," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 2818–2826.

[31] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," arXiv:1503.02531, 2015.

[32] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, "Quantization and training of neural networks for efficient integer-arithmetic-only inference," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 2704–2713.

[33] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2019.

[34] B. Efron, "Bootstrap methods: Another look at the jackknife," *The Annals of Statistics*, vol. 7, no. 1, pp. 1–26, 1979.

[35] Q. McNemar, "Note on the sampling error of the difference between correlated proportions or percentages," *Psychometrika*, vol. 12, no. 2, pp. 153–157, 1947.

[36] S. Holm, "A simple sequentially rejective multiple test procedure," *Scandinavian J. Statistics*, vol. 6, no. 2, pp. 65–70, 1979.